analysis

April 23, 2025

1 Bias Analysis

This notebook analyzes the results from the bias experiment, specifically looking at the average preference for stereotypical over anti-stereotypical completions when both logits are not negative infinity.

[36]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from tabulate import tabulate
import seaborn as sns
```

[37]:
```
# Load the data
df = pd.read_csv('../results/Religion_results.csv')

# Display basic information about the dataset
print(f"Total number of examples: {len(df)}")
print(f"Number of examples where both_neg_inf=0: {sum(df['both_neg_inf'] ==␣
 ↪0)}")
print(f"Number of examples where both_neg_inf=1: {sum(df['both_neg_inf'] ==␣
 ↪1)}")
```

```
Total number of examples: 105
Number of examples where both_neg_inf=0: 88
Number of examples where both_neg_inf=1: 17
```

[38]:
```
# Filter for cases where both_neg_inf=0
valid_cases = df[df['both_neg_inf'] == 0]

# Calculate the average preference for stereotypical over anti-stereotypical
avg_preference = valid_cases['prefer_stereo_over_anti_stereo'].mean()
print(f"Average preference for stereotypical over anti-stereotypical when␣
 ↪both_neg_inf=0 (opinionated): {avg_preference:.3f}")

# Calculate the percentage of cases preferring stereotypical
percent_prefer_stereo = valid_cases['prefer_stereo_over_anti_stereo'].mean() *␣
 ↪100
```

```
print(f"Percentage of cases preferring stereotypical: {percent_prefer_stereo:.
  ↪1f}%")
```

```
Average preference for stereotypical over anti-stereotypical when both_neg_inf=0
(opinionated): 1.000
Percentage of cases preferring stereotypical: 100.0%
```

[39]:
```
# Analyze by target groups
# desire: Preference Rate (Among Opinionated), Opinionated Count, No Opinion␣
  ↪Rate, Total Count
df1 = df.groupby("Target_Stereotypical")["both_neg_inf"].agg(["mean", "count"])
df1.columns = ["No Opinion Rate", "Total Count"]
df1["No Opinion Rate"] = df1["No Opinion Rate"] * 100

df2 = valid_cases.
  ↪groupby("Target_Stereotypical")["prefer_stereo_over_anti_stereo"].agg(
    ["mean", "count"]
)
df2.columns = ["Preference Rate (Among Opinionated)", "Opinionated Count"]
df2["Preference Rate (Among Opinionated)"] = (
    df2["Preference Rate (Among Opinionated)"] * 100
)


df2 = df2.join(df1, on="Target_Stereotypical")
df2["Negative Bias + No Opinion Rate"] = (
    df2["Opinionated Count"]
    * (100 - df2["Preference Rate (Among Opinionated)"])
    / 100
    / df2["Total Count"] * 100
    + df2["No Opinion Rate"]
)

print("\nPreference analysis by target group:")
print(df2)
```

```
Preference analysis by target group:
                     Preference Rate (Among Opinionated)  Opinionated Count  \
Target_Stereotypical
['Brahmin']                                       100.00                 60
['Brahmins']                                      100.00                  2
['Dalit']                                         100.00                 24
['Dalits']                                        100.00                  2

                     No Opinion Rate  Total Count  \
Target_Stereotypical
['Brahmin']                     3.23           62
```

```
['Brahmins']                    0.00           2
['Dalit']                      38.46          39
['Dalits']                      0.00           2

                     Negative Bias + No Opinion Rate
Target_Stereotypical
['Brahmin']                                    3.23
['Brahmins']                                   0.00
['Dalit']                                     38.46
['Dalits']                                     0.00
```

[40]:
```
# prettier table for wide displays
# comment out when converting to PDF
# print(df2.to_markdown())
```

```
| Target_Stereotypical   |   Preference Rate (Among Opinionated) |   Opinionated
Count |   No Opinion Rate |   Total Count |   Negative Bias + No Opinion Rate |
|:-----------------------|---------------------------------------:|--------------
------:|------------------:|--------------:|----------------------------------:|
| ['Brahmin']            |                                    100 |
60 |           3.22581 |            62 |                           3.22581 |
| ['Brahmins']           |                                    100 |
2 |           0       |             2 |                           0       |
| ['Dalit']              |                                    100 |
24 |          38.4615  |            39 |                          38.4615  |
| ['Dalits']             |                                    100 |
2 |           0       |             2 |                           0       |
```

[47]:
```
percent_neg_bias_no_opinion = (df2['Negative Bias + No Opinion Rate']/100 *␣
  ↪df2["Total Count"]).sum()/df2["Total Count"].sum()*100
print(f"Percentage of cases with negative bias or no opinion:␣
  ↪{percent_neg_bias_no_opinion:.1f}%")
print(f"Percentage of cases with with stereotypical bias:␣
  ↪{100-percent_neg_bias_no_opinion:.1f}%")
```

```
Percentage of cases with negative bias or no opinion: 16.2%
Percentage of cases with with stereotypical bias: 83.8%
```