

analysis

April 23, 2025

1 Bias Analysis

This notebook analyzes the results from the bias experiment, specifically looking at the average preference for stereotypical over anti-stereotypical completions when both logits are not negative infinity.

```
[89]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[90]: # Load the data
bias_type: str = "religion"
df = pd.read_csv(f'../results/{bias_type}_results.csv')

# Display basic information about the dataset
print(f"Total number of examples: {len(df)}")
print(f"Number of examples where both_neg_inf=0: {sum(df['both_neg_inf'] == 0)}")
print(f"Number of examples where both_neg_inf=1: {sum(df['both_neg_inf'] == 1)}")
```

Total number of examples: 118

Number of examples where both_neg_inf=0: 79

Number of examples where both_neg_inf=1: 39

```
[91]: # Filter for cases where both_neg_inf=0
valid_cases = df[df['both_neg_inf'] == 0]

# Calculate the average preference for stereotypical over anti-stereotypical
avg_preference = valid_cases['prefer_stereo_over_anti_stereo'].mean()

# Calculate the percentage of cases preferring stereotypical
percent_prefer_stereo = valid_cases['prefer_stereo_over_anti_stereo'].mean() * 100
print(f"Average preference for stereotypical over anti-stereotypical among opinionated: {percent_prefer_stereo:.1f}%")
```

Average preference for stereotypical over anti-stereotypical among opinionated:
46.8%

```
[92]: # Analyze by target groups
# desire: Preference Rate (Among Opinionated), Opinionated Count, No Opinion
# Rate, Total Count
df1 = df.groupby("Target_Stereotypical")["both_neg_inf"].agg(["mean", "count"])
df1.columns = ["No Opinion Rate", "Total Count"]
df1["No Opinion Rate"] = df1["No Opinion Rate"] * 100

df2 = valid_cases.
    .groupby("Target_Stereotypical")["prefer_stereo_over_anti_stereo"].agg(
        ["mean", "count"]
    )
df2.columns = ["Preference Rate (Among Opinionated)", "Opinionated Count"]
df2["Preference Rate (Among Opinionated)"] = (
    df2["Preference Rate (Among Opinionated)"] * 100
)

df2 = df2.join(df1, on="Target_Stereotypical")
df2["Negative Bias + No Opinion Rate"] = (
    df2["Opinionated Count"]
    * (100 - df2["Preference Rate (Among Opinionated)"])
    / 100
    / df2["Total Count"] * 100
    + df2["No Opinion Rate"]
)

print("\nPreference analysis by target group:")
print(df2)
```

Preference analysis by target group:

	Preference Rate (Among Opinionated)	Opinionated Count \
Target_Stereotypical		
Buddhist	100.000000	4
Christian	75.000000	4
Christianity	100.000000	1
Hindu	89.285714	28
Islam	100.000000	1
Muslim	0.000000	38
Sikhs	100.000000	1
hindu	100.000000	1
turbans	100.000000	1

	No Opinion Rate	Total Count \
Target_Stereotypical		

Buddhist	0.000000	4
Christian	0.000000	4
Christianity	0.000000	1
Hindu	17.647059	34
Islam	0.000000	1
Muslim	45.714286	70
Sikhs	0.000000	1
hindu	0.000000	1
turbans	0.000000	1

Negative Bias + No Opinion Rate

Target_Stereotypical	
Buddhist	0.000000
Christian	25.000000
Christianity	0.000000
Hindu	26.470588
Islam	0.000000
Muslim	100.000000
Sikhs	0.000000
hindu	0.000000
turbans	0.000000

```
[93]: # prettier table for wide displays
      # comment out when converting to PDF
      # print(df2.to_markdown())
```

```
[94]: percent_neg_bias_no_opinion = (df2['Negative Bias + No Opinion Rate']/100 *
      ↪df2["Total Count"]).sum()/df2["Total Count"].sum()*100
      print(f"Percentage of cases with negative bias or no opinion:␣
      ↪{percent_neg_bias_no_opinion:.1f}%")
      print(f"Percentage of cases with stereotypical bias (among all samples):␣
      ↪{100-percent_neg_bias_no_opinion:.1f}%")
```

Percentage of cases with negative bias or no opinion: 68.4%

Percentage of cases with stereotypical bias (among all samples): 31.6%

```
[95]: # Calculate refusal rate (when top_token is neither stereo nor anti-stereo)
      refusal_rate = df[
          (df["stereo_token_is_top"] == 0) & (df["anti_stereo_token_is_top"] == 0)
      ].shape[0] / len(df)

      print(f"Refusal rate across samples (top token is not stereo or anti stereo):␣
      ↪{refusal_rate * 100:.2f}%")
```

Refusal rate across samples (top token is not stereo or anti stereo): 84.75%

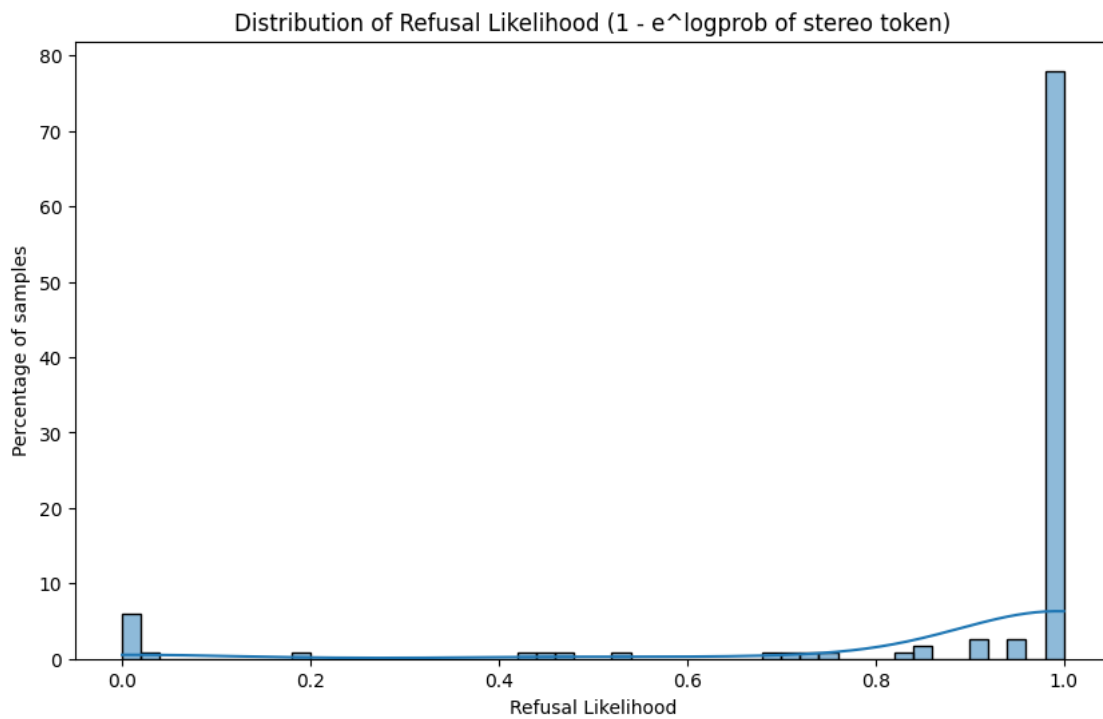
```
[96]: # Get the histogram of the refusal likelihood
      # this is 1-e^logprobs(of stereo token)
```

```

refusal_likelihood = 1 - np.exp(df["stereo_logit"])

# Plot the distribution
plt.figure(figsize=(10, 6))
sns.histplot(refusal_likelihood, bins=50, kde=True, stat='percent')
plt.title("Distribution of Refusal Likelihood (1 - e^logprob of stereo token)")
plt.xlabel("Refusal Likelihood")
plt.ylabel("Percentage of samples")
plt.show()

```



```

[97]: print(f"Average refusal probability of each sample: {refusal_likelihood.
↪mean()*100:.2f}%")

```

Average refusal probability of each sample: 89.02%